

# Лекция 4

- Уровни компиляции *nvcc*.
- *.cubin*, *.fatbin*, *.gpru* и *.ptx* файлы.
- PTX (*Parallel Thread eXecution*) ISA (*Instruction Set Architecture*).
- CUDA Driver API.

```
__global__ void gTest0(float* a, float* b){  
    int tid=threadIdx.x+blockIdx.x*blockDim.x;  
    a[tid]=a[tid]+b[tid];  
}  
int main(){  
.....  
    gTest0<<<N/128, 128>>>(a_d,b_d);  
    cudaDeviceSynchronize();  
.....  
}
```

```
tests/test0> nvcc -g -G test0m.cu -o test0m  
/tests/test0> cuda-gdb test0m
```

```
(cuda-gdb) list gTest0
```

```
1      #include <stdio.h>
2      #include <malloc.h>
3
4      __global__ void gTest0(float* a, float* b){
5          int tid=threadIdx.x+blockIdx.x*blockDim.x;
6          a[tid]=a[tid]+b[tid];
7      }
8
9      int main(){
10         int N=1024;
```

```
(cuda-gdb) break 5
```

```
..... •
```

(cuda-gdb) **run**

```
.....  
CUDA thread hit Breakpoint 1, gTest0<<<(8,1,1),(128,1,1)>>>  
(a=0x7fffc3200000, b=0x7fffc3201000)  
    at test0m.cu:5  
5          int tid=threadIdx.x+blockIdx.x*blockDim.x;
```

(cuda-gdb) **disassemble**

Dump of assembler code for function `_Z6gTest0PfS_`:

```
0x00007fffc8e3f800 <+0>:  MOV R1, c[0x0][0x28]
0x00007fffc8e3f810 <+16>:  MOV R2, RZ
0x00007fffc8e3f820 <+32>:  LDC.64 R2, c[0x0][R2+0x160]
0x00007fffc8e3f830 <+48>:  MOV R7, R2
0x00007fffc8e3f840 <+64>:  MOV R8, R3
0x00007fffc8e3f850 <+80>:  MOV R7, R7
0x00007fffc8e3f860 <+96>:  MOV R8, R8
```

```
.....
0x00007fffc8e3f8e0 <+224>:  MOV R8, R8
0x00007fffc8e3f8f0 <+240>:  MOV R5, R5
0x00007fffc8e3f900 <+256>:  MOV R6, R6
=> 0x00007fffc8e3f910 <+272>:  S2R R0, SR_TID.X
0x00007fffc8e3f920 <+288>:  MOV R0, R0
0x00007fffc8e3f930 <+304>:  S2R R2, SR_CTAID.X
0x00007fffc8e3f940 <+320>:  MOV R2, R2
```

```
.....
0x00007fffc8e3fca0 <+1184>:  BRA 0x4a0
0x00007fffc8e3fcb0 <+1200>:  NOP
0x00007fffc8e3fcc0 <+1216>:  NOP
```

# Раздельная компиляция

```
extern "C"  
__global__ void gTest2(float* a, float* b){  
    int tid=threadIdx.x+blockIdx.x*blockDim.x;  
    a[tid]=a[tid]+b[tid];  
}
```

***test2.cu***

```
#include <stdio.h>
#include <malloc.h>
```

***test2m.cu***

```
extern "C" __global__ void gTest2(float*, float*);
```

```
int main(){
    int N=1024;
    float* a=(float*)calloc(N, sizeof(float));
    float* b=(float*)calloc(N, sizeof(float));

    for(int i=0; i<N; i++){
        a[i]=2*i;
        b[i]=2*i+1;
    }
}
```

```
float *a_d, *b_d;
```

```
cudaMalloc((void**)&a_d, N*sizeof(float));
```

```
cudaMalloc((void**)&b_d, N*sizeof(float));
```

```
cudaMemcpy(a_d, a, N*sizeof(float), cudaMemcpyHostToDevice);
```

```
cudaMemcpy(b_d, b, N*sizeof(float), cudaMemcpyHostToDevice);
```

```
gTest2<<<N/128,128>>>(a_d,b_d);
```

```
cudaDeviceSynchronize();
```

```
cudaMemcpy(a, a_d, N*sizeof(float), cudaMemcpyDeviceToHost);
```

```
for(int i=0; i<N; i+=N/16)
```

```
printf("%g\n",a[i]);
```

```
}
```



```
tests/test2> nvcc -c test2.cu
```

```
tests/test2> nvcc -c test2m.cu
```

```
tests/test2> nvcc test2.o test2m.o -o test2e
```

```
#include <stdio.h>
#include <malloc.h>
#include <cuda.h>
```

```
extern "C" void hLauncherTest2(float* a, float* b, int N);
```

```
int main(){
```

```
.....
```

```
    hLauncherTest2(a_d,b_d,N);
```

```
.....
```

```
}
```

```
__global__ void gTest2(float* a, float* b){  
    int tid=threadIdx.x+blockIdx.x*blockDim.x;  
    a[tid]=a[tid]+b[tid];  
}
```

```
extern "C"  
void hLauncherTest2(float* a, float* b, int N){  
    gTest2<<<N/128,128>>>(a,b);  
    cudaDeviceSynchronize();  
}
```

```
nvcc -Xcompiler -fPIC -shared test2.cu -o libtest2.so
```

```
export LD_LIBRARY_PATH=$LD_LIBRARY_PATH:.
```

```
nvcc test2m.cu -L. -ltest2 -o test2m
```

# PTX

PTX (Parallel Thread eXecution) определяет виртуальную машину и набор инструкций (ISA - Instruction Set Architecture). PTX программа транслируется во время загрузки в команды соответствующего GPU и машинный код загружается на GPU драйвером.

```
tests/test2> nvcc -ptx test2.cu
```

```
.version 8.5  
.target sm_52  
.address_size 64
```

```
// .globl    gTest2
```

```
.visible .entry gTest2(  
    .param .u64 gTest2_param_0,  
    .param .u64 gTest2_param_1  
)  
{  
    .reg .f32    %f<4>;  
    .reg .b32    %r<5>;  
    .reg .b64    %rd<8>;
```

***test2.ptx***

```
ld.param.u64    %rd1, [gTest2_param_0];
ld.param.u64    %rd2, [gTest2_param_1];
cvta.to.global.u64    %rd3, %rd2;
cvta.to.global.u64    %rd4, %rd1;
mov.u32         %r1, %tid.x;
mov.u32         %r2, %ctaid.x;
mov.u32         %r3, %ntid.x;
mad.lo.s32      %r4, %r2, %r3, %r1;
mul.wide.s32    %rd5, %r4, 4;
add.s64         %rd6, %rd4, %rd5;
ld.global.f32   %f1, [%rd6];
add.s64         %rd7, %rd3, %rd5;
ld.global.f32   %f2, [%rd7];
add.f32         %f3, %f1, %f2;
st.global.f32   [%rd6], %f3;
ret;
```

```
}
```

```
#include <stdio.h>
#include <malloc.h>
```

*mst.cu*

```
extern "C" { __global__ void gStub(float* a, float* b){ } }
```

```
int main(){
    int N=2048;
    float* a=(float*)calloc(N, sizeof(float));
    float* b=(float*)calloc(N, sizeof(float));

    for(int i=0; i<N; i++){
        a[i]=2*i;
        b[i]=2*i+1;
    }
}
```



```
float *a_d, *b_d;
    cudaMalloc((void**)&a_d, N*sizeof(float));
    cudaMalloc((void**)&b_d, N*sizeof(float));

    cudaMemcpy(a_d, a, N*sizeof(float), cudaMemcpyHostToDevice);
    cudaMemcpy(b_d, b, N*sizeof(float), cudaMemcpyHostToDevice);

    gStub<<<N/128,128>>>(a_d,b_d);
    cudaDeviceSynchronize();

    cudaMemcpy(a, a_d, N*sizeof(float), cudaMemcpyDeviceToHost);

    for(int i=0; i<N; i+=N/16)
        printf("%g\n",a[i]);
}
```

```
.version 8.5
.target sm_52
.address_size 64
```

```
    // .globl    gStub
```

```
.visible .entry gStub(
    .param .u64 gStub_param_0,
    .param .u64 gStub_param_1
)
{

    ret;

}
```

```
tests/mystub> nvcc -dryrun -arch=sm_52 mst.cu -o mst  
--keep 2> dryrun_mst.out
```

```
tests/mystub> ll
```

```
total 12
```

```
-rw-r--r-- 1 malkov users 5053 abr 21 15:18
```

```
dryrun_mst.out
```

```
-rw-r--r-- 1 malkov users 815 abr 21 15:09 mst.cu
```

```
tests/mystub> ls -ltr
```

```
total 4260
```

```
drwxr-xr-x 2 malkov users    21 abr 21 13:42 spare
-rw-r--r-- 1 malkov users   815 abr 21 15:09 mst.cu
-rw-r--r-- 1 malkov users  5053 abr 21 15:18 dryrun_mst.out
-rw-r--r-- 1 malkov users 1091325 abr 21 15:21 mst.cpp4.ii
-rw-r--r-- 1 malkov users    24 abr 21 15:21 mst.module_id
-rw-r--r-- 1 malkov users 1010983 abr 21 15:21 mst.cudafe1.cpp
-rw-r--r-- 1 malkov users 1184988 abr 21 15:21 mst.cpp1.ii
-rw-r--r-- 1 malkov users    13 abr 21 15:21 mst.cudafe1.c
-rw-r--r-- 1 malkov users   312 abr 21 15:21 mst.ptx
-rw-r--r-- 1 malkov users   1237 abr 21 15:21 mst.cudafe1.stub.c
-rw-r--r-- 1 malkov users  10401 abr 21 15:21 mst.cudafe1.gpu
-rw-r--r-- 1 malkov users   1960 abr 21 15:21 mst.sm_52.cubin
-rw-r--r-- 1 malkov users   6647 abr 21 15:21 mst.fatbin.c
-rw-r--r-- 1 malkov users   2248 abr 21 15:21 mst.fatbin
-rw-r--r-- 1 malkov users   8984 abr 21 15:21 mst.o
-rw-r--r-- 1 malkov users    872 abr 21 15:21 mst_dlink.sm_52.cubin
-rw-r--r-- 1 malkov users    32 abr 21 15:21 mst_dlink.reg.c
-rw-r--r-- 1 malkov users   3190 abr 21 15:21 mst_dlink.fatbin.c
-rw-r--r-- 1 malkov users    952 abr 21 15:21 mst_dlink.fatbin
-rw-r--r-- 1 malkov users   2904 abr 21 15:21 mst_dlink.o
-rwxr-xr-x 1 malkov users  975856 abr 21 15:21 mst
```

```
nvcc mst.cu -keep -o mst
```

```
tests/mystub> ./mst
```

0

256

512

768

1024

1280

1536

1792

2048

2304

..... • •

```
extern "C"{  
__global__ void gStub(float* a, float* b){  
    int tid=threadIdx.x+blockIdx.x*blockDim.x;  
    a[tid]=a[tid]+b[tid];  
}  
}
```

```
.version 8.5
.target sm_52
.address_size 64
```

```
        // .globl          gStub
```

```
.visible .entry gStub(
    .param .u64 gStub_param_0,
    .param .u64 gStub_param_1
)
{
    .reg .f32      %f<4>;
    .reg .b32      %r<5>;
    .reg .b64      %rd<8>;
```

```
ld.param.u64      %rd1, [gStub_param_0];
ld.param.u64      %rd2, [gStub_param_1];
cvta.to.global.u64      %rd3, %rd2;
cvta.to.global.u64      %rd4, %rd1;
mov.u32           %r1, %tid.x;
mov.u32           %r2, %ctaid.x;
mov.u32           %r3, %ntid.x;
mad.lo.s32        %r4, %r2, %r3, %r1;
mul.wide.s32      %rd5, %r4, 4;
add.s64           %rd6, %rd4, %rd5;
ld.global.f32     %f1, [%rd6];
add.s64           %rd7, %rd3, %rd5;
ld.global.f32     %f2, [%rd7];
add.f32           %f3, %f1, %f2;
st.global.f32     [%rd6], %f3;
ret;
```

```
}
```



```
#$ cicc --c++17 --gnu_version=110200 --display_error_number --orig_src_file_name  
"mst.cu" --orig_src_path_name "/home/malkov/tests/mystub/mst.cu"  
--allow_managed -arch compute_52 -m64 --no-version-ident -ftz=0 -prec_div=1  
-prec_sqrt=1 -fmad=1 --include_file_name "mst.fatbin.c" -tused  
--module_id_file_name "mst.module_id" --gen_c_file_name "mst.cudafe1.c"  
--stub_file_name "mst.cudafe1.stub.c" --gen_device_file_name "mst.cudafe1.gpu"  
"mst.cpp1.ii" -o "mst.ptx"
```

```
#$-----  
ptxas -arch=sm_52 -m64 "mst.ptx" -o "mst.sm_52.cubin"
```

```
fatbinary --create="mst.fatbin" -64 --cicc-cmdline="-ftz=0 -prec_div=1 -prec_sqrt=1  
-fmad=1 " "--image3=kind=elf,sm=52,file=mst.sm_52.cubin"  
"--image3=kind=ptx,sm=52,file=mst.ptx" --embedded-fatbin="mst.fatbin.c"
```

```
"/home/malkov/anaconda3/bin"/x86_64-conda-linux-gnu-c++  
-D__CUDA_ARCH__=520 -D__CUDA_ARCH_LIST__=520  
-D__NV_LEGACY_LAUNCH -c -x c++ -DCUDA_DOUBLE_MATH_FUNCTIONS  
-Wno-psabi "-I/usr/local/cuda-12.5/bin/./targets/x86_64-linux/include" -m64  
"mst.cudafe1.cpp" -o "mst.o"
```

```
nvlink -m64 --arch=sm_52 --register-link-binaries="mst_dlink.reg.c"  
"-L/usr/local/cuda-12.5/bin/./targets/x86_64-linux/lib/stubs"  
"-L/usr/local/cuda-12.5/bin/./targets/x86_64-linux/lib" -cpu-arch=X86_64 "mst.o"  
-lcudadevrt -o "mst_dlink.sm_52.cubin" --host-ccbin  
"/home/malkov/anaconda3/bin/x86_64-conda-linux-gnu-c++"
```

```
fatbinary --create="mst_dlink.fatbin" -64 --cicc-cmdline="-ftz=0 -prec_div=1  
-prec_sqrt=1 -fmad=1 " -link "--image3=kind=elf,sm=52,file=mst_dlink.sm_52.cubin"  
--embedded-fatbin="mst_dlink.fatbin.c"
```

```
"/home/malkov/anaconda3/bin"/x86_64-conda-linux-gnu-c++  
-D__CUDA_ARCH_LIST__=520 -D__NV_LEGACY_LAUNCH -c -x c++  
-DFATBINFILE="\"mst_dlink.fatbin.c\""  
-DREGISTERLINKBINARYFILE="\"mst_dlink.reg.c\"" -l.  
-D__NV_EXTRA_INITIALIZATION= -D__NV_EXTRA_FINALIZATION=  
-D__CUDA_INCLUDE_COMPILER_INTERNAL_HEADERS__ -Wno-psabi  
"-l/usr/local/cuda-12.5/bin/./targets/x86_64-linux/include"  
-D__CUDACC_VER_MAJOR__=12 -D__CUDACC_VER_MINOR__=5  
-D__CUDACC_VER_BUILD__=82 -D__CUDA_API_VER_MAJOR__=12  
-D__CUDA_API_VER_MINOR__=5 -D__NVCC_DIAG_PRAGMA_SUPPORT__=1  
-m64 "/usr/local/cuda-12.5/bin/crt/link.stub" -o "mst_dlink.o"
```

```
"/home/malkov/anaconda3/bin"/x86_64-conda-linux-gnu-c++  
-D__CUDA_ARCH_LIST__=520 -D__NV_LEGACY_LAUNCH -m64 -Wl,--start-group  
"mst_dlink.o" "mst.o" "-L/usr/local/cuda-12.5/bin/./targets/x86_64-linux/lib/stubs"  
"-L/usr/local/cuda-12.5/bin/./targets/x86_64-linux/lib" -lcudadevrt -lcudart_static -lrt  
-lpthread -ldl -Wl,--end-group -o "mst"
```

```
tests/mystub> ./dryrun_mst.out
```

```
tests/mystub> ./mst
```

```
1
```

```
513
```

```
1025
```

```
1537
```

```
2049
```

```
2561
```

```
3073
```

```
3585
```

```
4097
```

```
4609
```

```
..... • •
```



```
tests/cudrapi> g++ -I/usr/local/cuda/include  
-L/usr/local/cuda/lib64 -lcudart -lcuda cda.cpp -o cda
```

```
#include <cuda.h>
#include <cuda_runtime.h>
#include <stdio.h>
#include <malloc.h>
```

```
int main(){
    cuInit(0);
    CUdevice cuDevice;
    CUresult res = cuDeviceGet(&cuDevice, 0);
    if (res != CUDA_SUCCESS){
        printf("cannot acquire device 0\n");
        exit(1);
    }

    CUcontext cuContext;
    res = cuCtxCreate(&cuContext, 0, cuDevice);
    if (res != CUDA_SUCCESS){
        printf("cannot create context\n");
        exit(1);
    }
}
```



```
int N=2048;  
float* a=(float*)calloc(N, sizeof(float));  
float* b=(float*)calloc(N, sizeof(float));
```

```
for(int i=0; i<N; i++){  
    a[i]=2*i;  
    b[i]=2*i+1;  
}
```

```
float *a_d, *b_d;  
cudaMalloc((void**)&a_d, N*sizeof(float));  
cudaMalloc((void**)&b_d, N*sizeof(float));
```

```
cudaMemcpy(a_d, a, N*sizeof(float), cudaMemcpyHostToDevice);  
cudaMemcpy(b_d, b, N*sizeof(float), cudaMemcpyHostToDevice);
```

```
//gStub<<<N/128,128>>>(a_d,b_d);  
//cudaDeviceSynchronize();
```

```
CUmodule cuModule = (CUmodule)0;
cuModuleLoad(&cuModule, "cda.ptx");
CUfunction gStub;
cuModuleGetFunction(&gStub, cuModule, "gStub");

void* args[] = {&a_d, &b_d};
cuLaunchKernel(gStub, N/128, 1, 1, 128, 1, 1, 0, 0, args, 0);

cudaMemcpy(a, a_d, N*sizeof(float), cudaMemcpyDeviceToHost);

for(int i=0; i<N; i+=N/16)
    printf("%g\n",a[i]);

    cuCtxDestroy(cuContext);
return 0;
}
```